

resnext32_pretrained model report:

Hyperparameters:

Epochs: 10

Max lr: 0.001

Optimizer: Adam

Loss function: cross entropy

Training log:

Dataset has 102 examples and 2 classes.

Training size: 72, Test size: 30.

100%|██████████| 3/3 [00:10<00:00, 3.40s/it]

Epoch [0] : train_loss: 0.7026, train_acc: 0.4688

100%|██████████| 3/3 [00:10<00:00, 3.51s/it]

Epoch [1] : train_loss: 0.7046, train_acc: 0.6250

100%|██████████| 3/3 [00:10<00:00, 3.41s/it]

Epoch [2] : train_loss: 0.6460, train_acc: 0.6979

100%|██████████| 3/3 [00:11<00:00, 3.72s/it]

Epoch [3] : train_loss: 0.5955, train_acc: 0.5312

100%|██████████| 3/3 [00:10<00:00, 3.47s/it]

Epoch [4] : train_loss: 0.6275, train_acc: 0.6875

100%|██████████| 3/3 [00:10<00:00, 3.43s/it]

Epoch [5] : train_loss: 0.5830, train_acc: 0.7917

100%|██████████| 3/3 [00:10<00:00, 3.48s/it]

Epoch [6] : train_loss: 0.6219, train_acc: 0.8333

100%|██████████| 3/3 [00:10<00:00, 3.56s/it]

Epoch [7] : train_loss: 0.5565, train_acc: 0.8229

100%|██████████| 3/3 [00:10<00:00, 3.37s/it]

Epoch [8] : train_loss: 0.5967, train_acc: 0.7500

100%|██████████| 3/3 [00:10<00:00, 3.34s/it]

Epoch [9] : train_loss: 0.5584, train_acc: 0.6250

Test accuracy: {'loss': 0.502370297908783, 'acc': 0.8333333134651184}.

Training accuracy curve and loss curve:

